

Decomposition of Drug Substance Use on Health Dispositions

AJ Marcus

2025-02-27

Preamble

Modules

```
In [1]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import pyreadr
from random import random
from scipy.stats import chi2_contingency, skewnorm
import seaborn as sns
from sklearn.compose import ColumnTransformer
from sklearn.decomposition import PCA
from sklearn.ensemble import (
    GradientBoostingRegressor,
    RandomForestClassifier,
    RandomForestRegressor
)
from sklearn.feature_selection import SequentialFeatureSelector
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import (
    accuracy_score,
    classification_report,
    mean_absolute_error,
    mean_squared_error,
    r2_score
)
from sklearn.preprocessing import (
    LabelEncoder,
    MinMaxScaler,
    OneHotEncoder,
    OrdinalEncoder,
    PowerTransformer,
    QuantileTransformer
)
from statsmodels.stats.outliers_influence import variance_inflation_factor
```

Functions

```
In [2]: def restrata_cat(cpd):
        i = random()
        for idx, x in cpd.items():
            if i > x:
                continue
            else:
                return idx
```

```
In [3]: def restrata_num(mean, std, bounds = None, skew = 0):
        while True:
            x = skewnorm.rvs(a= skew, size= 1)[0]
            x = (x * std) + mean
            if bounds:
                if (x >= bounds[0]) and (x <= bounds[1]):
                    return x
                else:
                    continue
            else:
                return x
```

Imports

```
In [4]: inread = pyreadr.read_r(".data/NSDUH_2023.Rdata")
```

Mining & Preprocessing

Subset features to information pertinent to research and analysis hypotheses. Research focuses on physical & mental health quality, drug dependencies, and recent drug uses. Most features utilized are recoded versions of original survey questions, as recoded by survey providers to create consistent, reliable, and viable data points.

Several features are not applicable to minors, which sources the primary reason for excluding them from final analysis. Segmentation analysis by respondent generation will be used to evaluate any necessary grouping or faceting in the report's findings.

```
In [5]: subset = [
        # DEMOGRAPHICS
        # Variable                                Name                                Data Type    Pre
        # -----
        'AGE3',                                   # age                                cat int
        'IRSEX',                                   # gender                             cat int
        'SEXIDENT22',                             # sexual identity                    cat int
        'IRMARIT',                                 # marital status                     cat int
        'IREDUHIGHST2',                           # education level                    ord int
        'EDUSCHLGO',                              # in-school status                   cat int
        'IRWRKSTAT',                              # employment status                  cat int
        # 'MILSTAT',                              # military status                    cat int
        'ANYHLT12',                                # healthcare status                  cat int
        'GOVTPROG',                               # income assistance status           cat int
        'INCOME',                                 # annual income                      ord int
```

```

    'POVERTY3',          # poverty status          ord int

# HEALTH
# Variable              Name                      Data Type
# -----
'HEALTH',              # health condition          ord int
'BMI2',                # BMI                      float 9-69
'KSSLR6MON',           # psych distress score      int 1-24
'AMIPY',               # any mental illness -yr    bool

# DRUG USE
# Variable              Name                      Data Type      Ris
# -----
'IRPMNICDEP',          # nicotine dependence -yr   bool
'IRMJFM',              # pot freq use -mth         int 1-30        1
'IRCOCFM',             # cocaine freq use -mth     int 1-30        6
'IRCRKFM',             # crack freq use -mth       int 1-30        6
'IRHERFM',             # heroine freq use -mth     int 1-30        7
'IRHALLUC30N',         # hallucinogens freq use -mth int 1-30        2
'IRINHAL30N',          # inhalants freq use -mth   int 1-30        3
'IRMETHAM30N',         # meth freq use -mth        int 1-30        5
'IRPNRNM30FQ',         # painmed freq use -mth     int 1-30        8
'IRTRQNM30FQ',         # tranqs freq use -mth      int 1-30        4
'IRSTMNM30FQ',         # stims freq use -mth       int 1-30        5
'IRSEDNM30FQ',         # sedatives freq use -mth   int 1-30        4
]

data = inread['puf2023_102124'][subset]
data = data[data['AGE3'] > 3]

```

Relabel features for readability.

```

In [6]: renames = {
    'AGE3': 'generation',
    'ANYHLTI2': 'health_insurance_flag',
    'BMI2': 'bmi',
    'EDUSCHLGO': 'education_flag',
    'GOVTPROG': 'income_assistance_flag',
    'HEALTH': 'health_level',
    'IRCOCFM': 'cocaine_use',
    'INCOME': 'annual_income',
    'IRCRKFM': 'crack_use',
    'IREDUHIGHEST2': 'education_level',
    'IRHALLUC30N': 'hallucinogen_use',
    'IRHERFM': 'heroine_use',
    'IRINHAL30N': 'inhalant_use',
    'IRMARIT': 'marital',
    'IRMETHAM30N': 'amphetamine_use',
    'IRMJFM': 'marijuana_use',
    'IRPMNICDEP': 'nicotine_dependence_flag',
    'IRPNRNM30FQ': 'opioid_use',
    'IRSEDNM30FQ': 'sedative_use',
    'IRSEX': 'gender',
    'IRSTMNM30FQ': 'stimulant_use',
    'IRWRKSTAT': 'employment',

```

```

    'IRTRQNM30FQ': 'tranquilizer_use',
    'KSSLR6MON': 'psychological_distress',
#    'MILSTAT': 'military_service',
    'POVERTY3': 'poverty_level',
    'SEXIDENT22': 'sexuality',
    'AMIPY': 'mental_illness_flag'
}

data = data.rename(columns= renames)

```

Recoding

Recoding variables from system codes to readable elements. Some overwriting recoding has been done to retain usability of some features, and is annotated accordingly.

Recode Edits:

1. Annual Income was grouped into survey ranges; groups are being populated to maintain uniform distributions within each band to allow smoother normalization in model preprocessing efforts.
2. Slight deviations of 0-3 years were made when switching age buckets to generations and ages were grouped into the generation which majority-mapped to the applicable range.

```

In [7]: # recoding
recodes = {
    'education_flag': {
        1: 'True',
        2: 'False',
        11: 'True'
    },
    'education_level': {          # effective US grade level
        1: 5,
        2: 6,
        3: 7,
        4: 8,
        5: 9,
        6: 10,
        7: 11,
        8: 12,
        9: 12,
        10: 13,
        11: 14
    },
    'employment': {
        1: 'Full Employment',
        2: 'Partial Employment',
        3: 'Unemployed',
        4: pd.NA,
        99: pd.NA
    },
    'gender': {

```

```

    1: 'Male',
    2: 'Female'
},
'generation': {                                # see recode note 2
    1: 'A',
    2: 'Z',
    3: 'Zillennial',
    4: 'Zillennial',
    5: 'Zillennial',
    6: 'Zillennial',
    7: 'Millennial',
    8: 'Millennial',
    9: 'Xennial',
    10: 'X',
    11: 'Baby Boomer'
},
'health_insurance_flag': {
    1: 'True',
    2: 'False'
},
'health_level': {
    1: 5,
    2: 4,
    3: 3,
    4: 2,
    5: 1
},
'income_assistance_flag': {
    1: 'True',
    2: 'False'
},
'marital': {
    None: 'Single',
    1: 'Married',
    2: 'Widowed',
    3: 'Divorced',
    4: 'Single'
},
'nicotine_dependence_flag': {
    0: 'False',
    1: 'True'
},
'poverty_level': {
    1: 'Below',
    2: 'Baseline',
    3: 'Above'
},
'sexuality': {
    1: 'Heterosexual',
    2: 'Homosexual',
    3: 'Bisexual'
},
'mental_illness_flag': {
    0: 'False',
    1: 'True'
}

```

```

    }
}

for c in recodes:
    mapper = recodes[c]
    data[c] = data[c].map(mapper)

```

Drug Use fields have 90-series codes referring to insufficient, bad, or non-applicable data. Features with such values are recoded to no use for conservative utility of the fields and biases against the hypothesis to increase the threshold of positive analysis results.

```

In [8]: reranges = [
    'amphetamine_use',
    'cocaine_use',
    'crack_use',
    'hallucinogen_use',
    'heroin_use',
    'inhalant_use',
    'marijuana_use',
    'opioid_use',
    'sedative_use',
    'stimulant_use',
    'tranquilizer_use'
]

for c in reranges:
    data[c] = data[c].apply(lambda x: 0 if x > 30 else x)

```

Clean datatypes

```

In [9]: data['education_level'] = data['education_level'].astype(float)
data['annual_income'] = data['annual_income'].astype(float)

```

Missing at Random Analysis

Several features show missing data. Missing at Random (MAR) analyses are done to reduce inter-dependency analysis and feature size for the model selection stages. Dependency is tested via Chi-Squared tests at the $\alpha = 0.05$ level.

```

In [10]: data.isnull().sum()

```

```

Out[10]: generation          0
gender                      0
sexuality                   4956
marital                     0
education_level             0
education_flag              1688
employment                  13444
health_insurance_flag       1740
income_assistance_flag      0
annual_income               0
poverty_level               11
health_level                14
bmi                         1964
psychological_distress      0
mental_illness_flag         0
nicotine_dependence_flag    0
marajuana_use               0
cocaine_use                 0
crack_use                   0
heroine_use                 0
hallucinogen_use           0
inhalant_use                0
amphetamine_use             0
opioid_use                  0
tranquilizer_use            0
stimulant_use               0
sedative_use                0
dtype: int64

```

```

In [11]: data_missing = data.isna()
colset = []

for i in range(len(data_missing.columns)):
    for j in range(i, len(data_missing.columns)):
        combo = sorted([data_missing.columns[i], data_missing.columns[j]])
        if combo[0] == combo[1]:
            continue
        if combo not in colset:
            colset.append(combo)

missing_tests = []
for a, b in colset:
    contingency = pd.crosstab(data_missing[a], data_missing[b])
    chi2, p, dof, expected = chi2_contingency(contingency)
    if p < .05:
        result = 'Dependent'
    else:
        result = 'Independent'
    if dof > 0:
        missing_tests.append(
            (a, b, round(p, 4), result)
        )

cross_missing = pd.DataFrame(missing_tests, columns= ['feature_a', 'feature_b', 'p_
cross_missing = cross_missing[cross_missing['dependency'] == 'Dependent']

```

```
cross_missing = cross_missing.sort_values(['p_value', 'feature_a', 'feature_b']).re
cross_missing
```

Out[11]:

	feature_a	feature_b	p_value	dependency
0	bmi	education_flag	0.0000	Dependent
1	bmi	employment	0.0000	Dependent
2	bmi	health_insurance_flag	0.0000	Dependent
3	bmi	health_level	0.0000	Dependent
4	bmi	sexuality	0.0000	Dependent
5	education_flag	employment	0.0000	Dependent
6	education_flag	health_insurance_flag	0.0000	Dependent
7	education_flag	health_level	0.0000	Dependent
8	education_flag	sexuality	0.0000	Dependent
9	employment	sexuality	0.0000	Dependent
10	health_insurance_flag	health_level	0.0000	Dependent
11	health_insurance_flag	sexuality	0.0000	Dependent
12	health_level	sexuality	0.0000	Dependent
13	employment	health_insurance_flag	0.0316	Dependent

We can see the dependency tests show 7 features dependent on each other in some combination. Preliminary EDA revealed other features that were revoked in report refinement (i.e. `military_service`) and others have a real world relationship we expect some dependency on (i.e. `employment` and `health_insurance_flag`). Results indicate the majority of these nulls can be restratified with appropriate filling scalers and transformations, while the `employment` and `health_insurance_flag` pair may not be applicable due to them not assuredly be MAR given the confidence bounds of the test.

Null Cleanup & Restrata

This is done only on previously-identified applicable fields (or on fields that did not result in the MAR analysis) that can be reasonably expected to not impact clustering or predictive analysis and is not done on fields with high null content.

```
In [12]: restratas_cat = [
    'education_flag',
    'health_insurance_flag',
    'health_level',
    'psychological_distress',
    'sexuality'
]
```



```

for c in restratas_cat:
    a = data[c].value_counts(normalize= True)
    cpd = a.sort_index().cumsum()
    data[c] = data[c].apply(lambda x: x if pd.notna(x) else restrata_cat(cpd))

```

```

In [13]: restratas_num = [
        'bmi'
    ]

for c in restratas_num:
    a, b = data[c].min(), data[c].max()
    m, d, s = data[c].mean(), data[c].std(), data[c].skew()
    data[c] = data[c].apply(lambda x: x if pd.notna(x) else restrata_num(m, d, boun

```

All remaining nulls are dropped to create the final cleaned dataset for modeling purposes.

```

In [14]: data_clean = data.dropna()
        data_clean.shape

```

Out[14]: (31683, 27)

Distribution Analysis

```



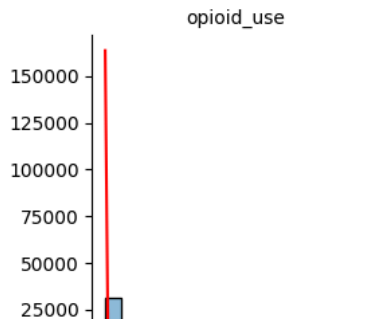
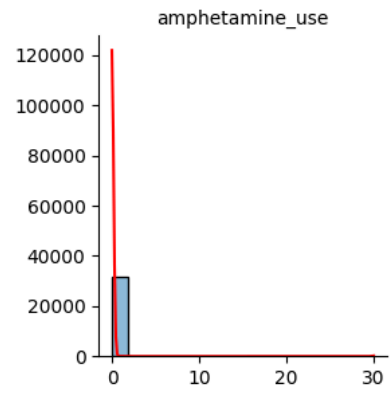
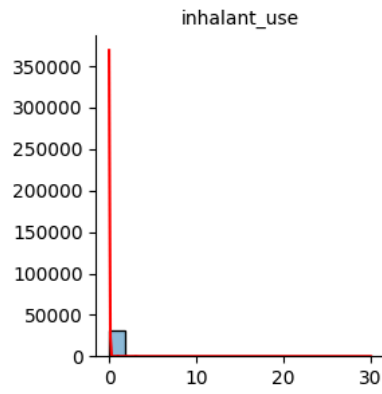

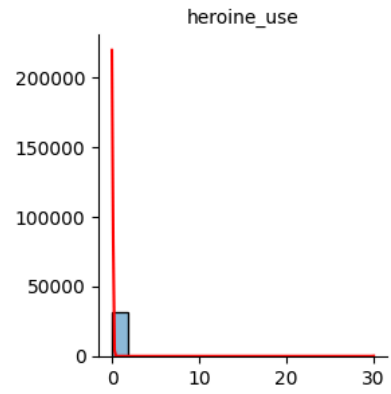
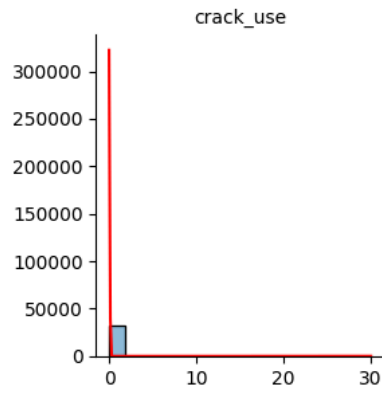
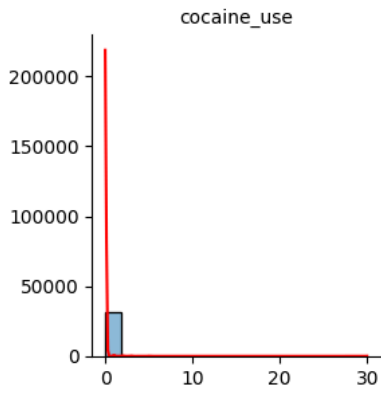
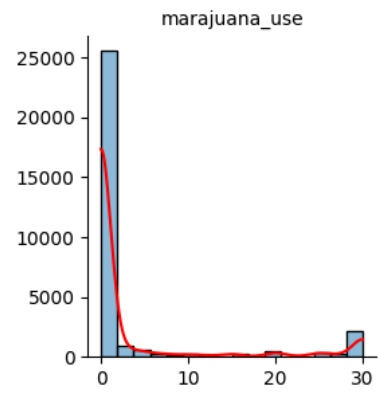
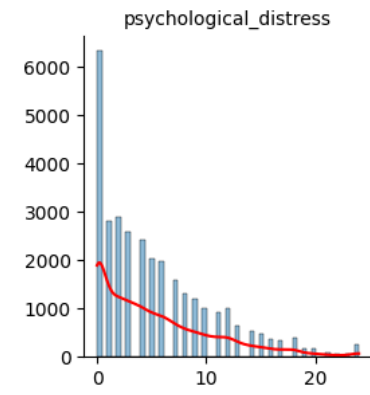

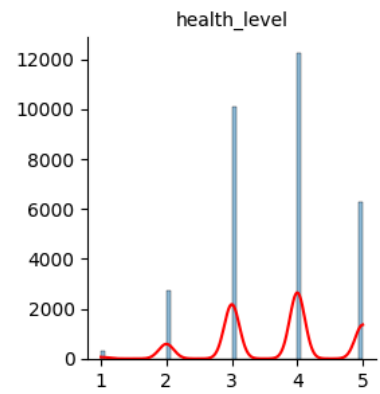
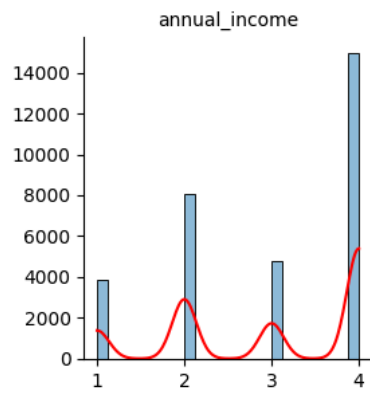
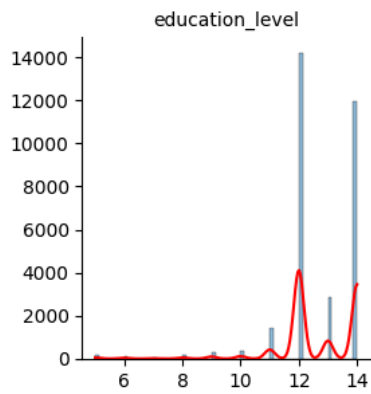
In [15]: cols_nums = data_clean.select_dtypes(include= 'number').columns.tolist()
        df_melt = data_clean[cols_nums].melt(
            var_name= 'col',
            value_name='val'
        )


        g = sns.FacetGrid(
            df_melt,
            col= 'col',
            col_wrap= 3,
            sharex= False,
            sharey= False
        )
        g.map(
            sns.histplot,
            'val',
            kde= True
        ).add_legend()

        for ax in g.axes.flat:
            for line in ax.lines:
                line.set_color('red')

        g.set_axis_labels('', '')
        g.set_titles("{col_name}")
        g.tight_layout()
        plt.show()

```



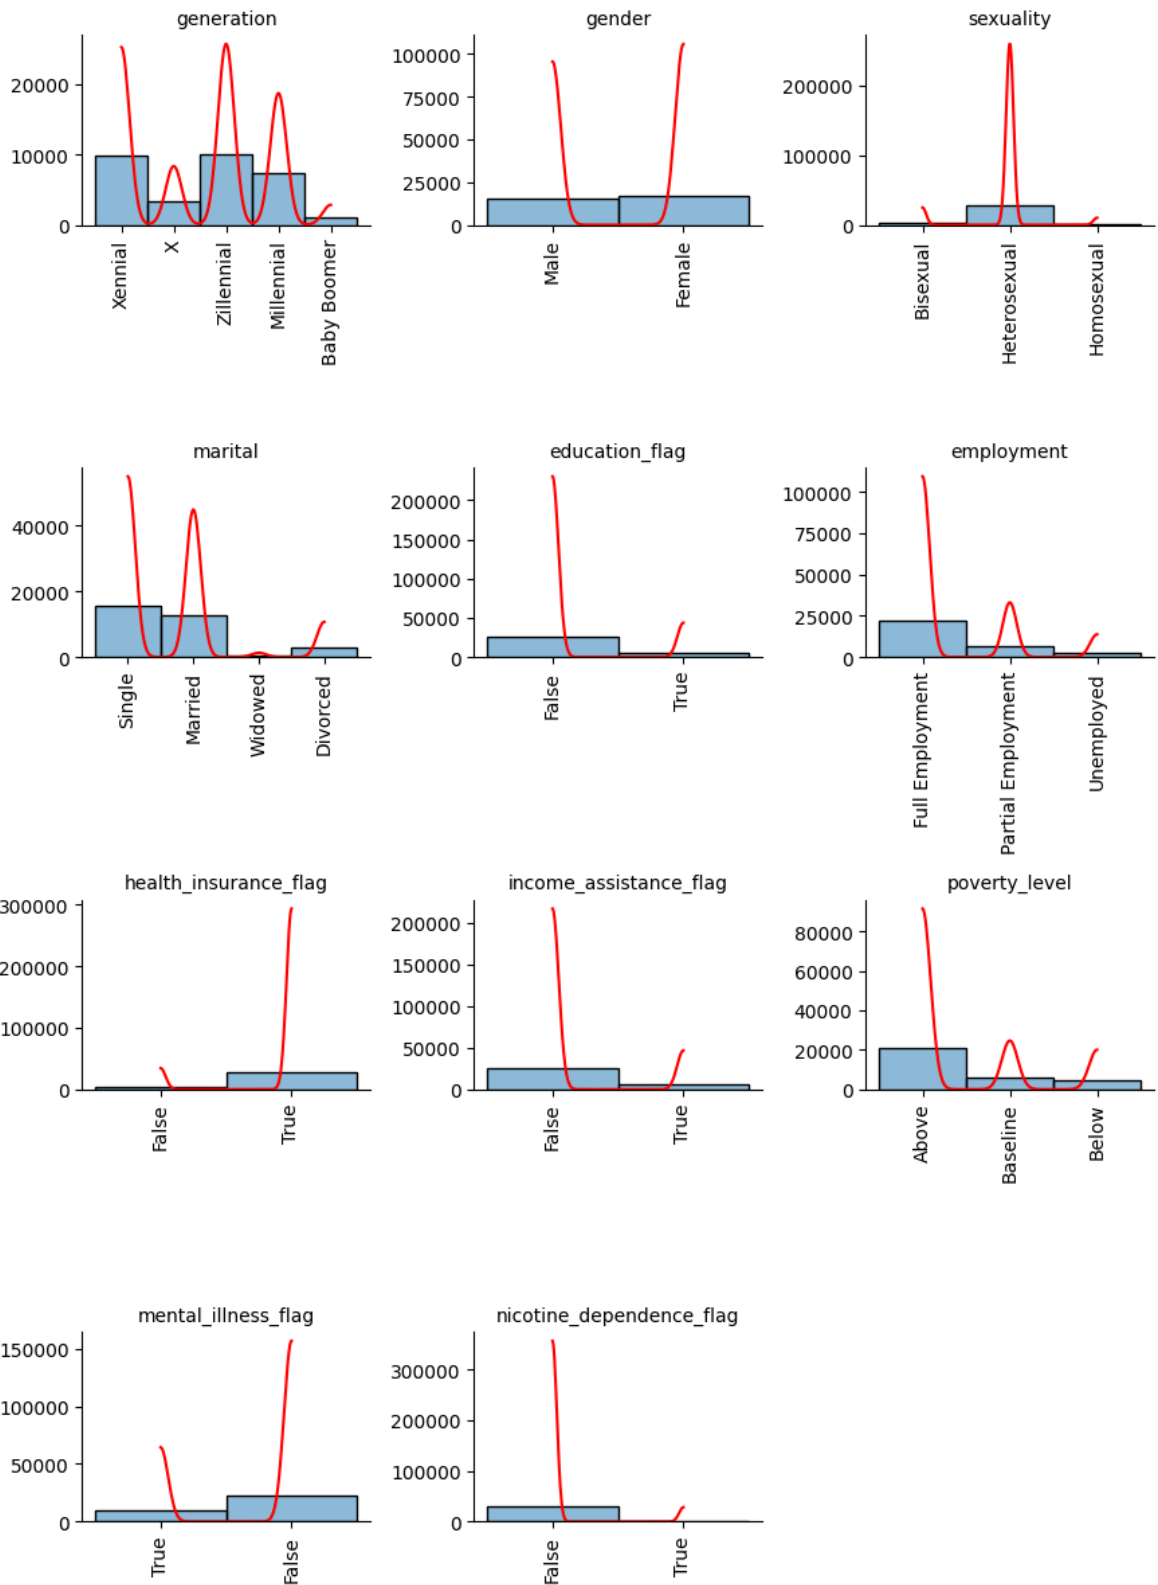


```
In [16]: cols_cats = data_clean.select_dtypes(exclude= 'number').columns.tolist()
df_melt = data_clean[cols_cats].melt(
    var_name= 'col',
    value_name='val'
)

g = sns.FacetGrid(
    df_melt,
    col='col',
    col_wrap= 3,
    sharex= False,
    sharey= False
)
g.map(
    sns.histplot,
    'val',
    kde= True
).add_legend()

for ax in g.axes.flat:
    for label in ax.get_xticklabels():
        label.set_rotation(90)
    for line in ax.lines:
        line.set_color('red')

# Adjust layout
g.set_axis_labels('', '')
g.set_titles("{col_name}")
g.tight_layout()
plt.show()
```



Scaling and Preprocessing

Distribution analysis shows the majority of features are not applicable to standard transformations for modeling. Drug use fields were chosen to apply

PowerTransformation in order to handle the heavily skewed data creating outlier and

distortion. Each modeling section creates its transformations applicable to the model subsets.

Performance was evaluated in some model applications on the transformations used in lieu of `PowerTransformation` with negligible-to-marginally negative results over the selected distribution. Those results are not shown here for concision.

```
In [17]: cols_minmax = [  
    'bmi'  
]  
  
cols_onehot = [  
    'education_flag',  
    'employment',  
    'gender',  
    'health_insurance_flag',  
    'income_assistance_flag',  
    'marital',  
    'mental_illness_flag',  
    'nicotine_dependence_flag',  
    'poverty_level',  
    'sexuality'  
]  
  
cols_ordinal = [  
    'generation'  
]  
  
cols_poweryeo = [  
    'amphetamine_use',  
    'cocaine_use',  
    'crack_use',  
    'hallucinogen_use',  
    'heroin_use',  
    'inhalant_use',  
    'marijuana_use',  
    'opioid_use',  
    'psychological_distress',  
    'sedative_use',  
    'stimulant_use',  
    'tranquilizer_use'  
]  
  
cols_quantnorm = [  
    'education_level',  
    'health_level'  
]  
  
cols_quantuni = [  
    'annual_income'  
]
```

Principal Component Analysis

The goal of this PCA is to try and identify key features that heavily-classify the data. Outputs from the PCA will be used to see if the highest-classifying features can be used to accurately and simply classify the data in other models.

```
In [18]: # preprocess transformations
tf_minmax = MinMaxScaler()
tf_onehot = OneHotEncoder()
tf_ordinal = OrdinalEncoder()
tf_poweryeo = PowerTransformer(method= 'yeo-johnson')
tf_quantnorm = QuantileTransformer(output_distribution= 'normal')
tf_quantuni = QuantileTransformer(output_distribution= 'uniform')

preprocessor = ColumnTransformer(
    transformers=[
        ('tf1', tf_minmax, cols_minmax),
        ('tf2', tf_onehot, cols_onehot),
        ('tf3', tf_ordinal, cols_ordinal),
        ('tf5', tf_poweryeo, cols_poweryeo),
        ('tf6', tf_quantnorm, cols_quantnorm),
        ('tf7', tf_quantuni, cols_quantuni)
    ]
)

data_processed = preprocessor.fit_transform(data_clean)

# explained variances
evr = []
for n in range(3, len(preprocessor.get_feature_names_out())):
    model_pca = PCA(n_components= n)
    pca_fit = model_pca.fit(data_processed)
    evr.append((n, sum(model_pca.explained_variance_ratio_)))
```

```
In [19]: z = []
for i, (x, y) in enumerate(evr):
    if i == len(evr) - 1:
        continue
    z.append(evr[i + 1][1] - y)

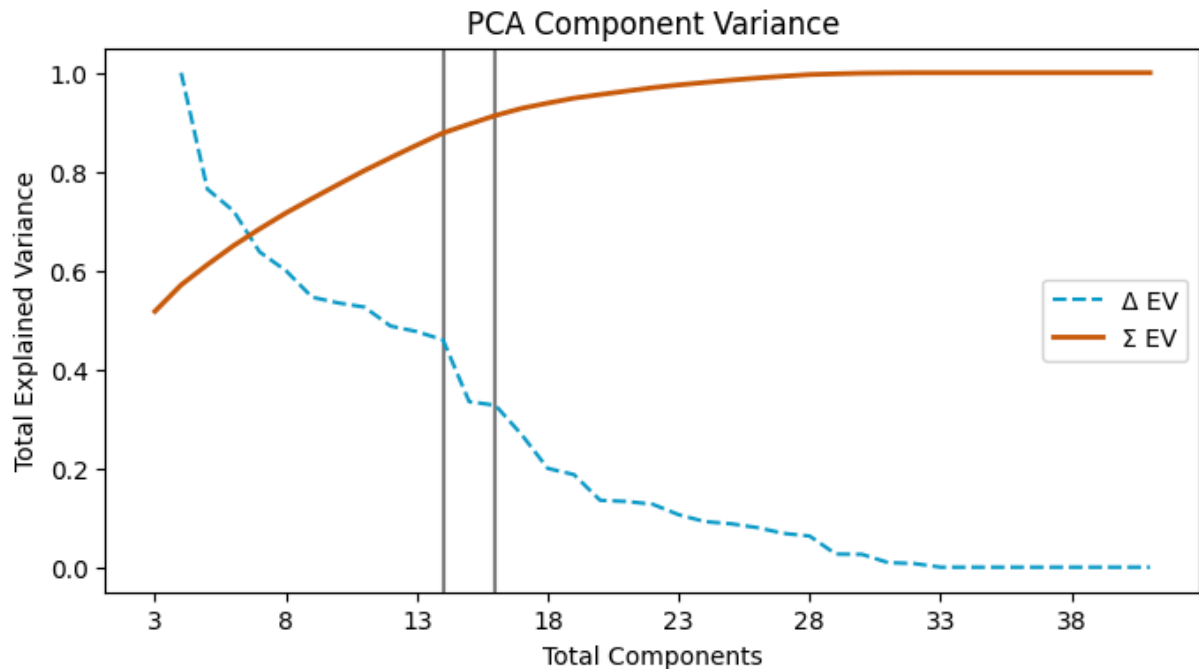
z = z / max(z)
x, y = zip(*evr)

plt.figure(figsize= (8, 4))
plt.axvline(x= 14, color= 'gray')
plt.axvline(x= 16, color= 'gray')
plt.plot(
    x[1:], z,
    linestyle= '--',
    color= '#0c9bc9',
    label= 'Δ EV'
)
plt.plot(
    x, y,
    linestyle= '-',
    color= '#c95b0c',
```

```

        linewidth= 2,
        label= ' $\Sigma$  EV'
    )
    plt.xticks(np.arange(3, max(x) + 1, 5))
    plt.xlabel('Total Components')
    plt.ylabel('Total Explained Variance')
    plt.title('PCA Component Variance')
    plt.legend(loc= 'center right')
    plt.show()

```



```

In [20]: oc = [14, 16]
        for i in oc:
            model_pca = PCA(n_components= i)
            pca_fit = model_pca.fit(data_processed)
            oev = sum(model_pca.explained_variance_ratio_)
            print(f"Optimal Explained Variance of {oev:.4f} with {i} Components")

```

Optimal Explained Variance of 0.8782 with 14 Components
 Optimal Explained Variance of 0.9136 with 16 Components

```

In [21]: pca_feature_names = preprocessor.get_feature_names_out()
        pca_top_components = model_pca.components_[:max(oc)]
        pca_components = pd.DataFrame()

        for i, component in enumerate(pca_top_components):
            sorted_feature_idx = np.argsort(np.abs(component))[:, -1]
            sorted_feature_names = preprocessor.get_feature_names_out()[sorted_feature_idx]
            sorted_loadings = component[sorted_feature_idx]

            row = dict(zip(sorted_feature_names, sorted_loadings))
            pca_components = pd.concat(
                [pca_components, pd.DataFrame([row])],
                ignore_index= True
            )

```

Borda Influence Scores

Results showed 16 component vectors the optimal total variance coverage. To determine which features have the highest influence over the components, we will use weighted borda summations where each subsequent component vector has a linearly devalued weight and the borda score values of each vector are calculated. The lower the weighted score, the more value it contributed to the explained variance within each component vector.

```
In [22]: pca_compranks = pca_components.T.abs().rank(method= 'dense', ascending= False)
for i in pca_compranks.columns:
    pca_compranks[i] = pca_compranks[i] * (16 - i)
borda_scores = round(pca_compranks.sum(axis= 1) / 16, 1)
borda_scores.sort_values(ascending= True)[:10]
```

```
Out[22]: tf5__marajuana_use          58.3
tf3__generation                    65.9
tf5__psychological_distress        69.1
tf5__amphetamine_use              79.6
tf5__cocaine_use                   84.6
tf5__opioid_use                    85.8
tf5__heroin_use                    92.2
tf5__tranquilizer_use             101.4
tf5__crack_use                    101.8
tf5__hallucinogen_use             103.7
dtype: float64
```

Multicollinearity

The downside is a potential concern of multicollinearity or intra-dependence, which will be evaluated next utilizing a variance inflation factor analysis.

```
In [23]: mc_cols = [
    'psychological_distress',
    'marajuana_use',
    'hallucinogen_use',
    'cocaine_use',
    'opioid_use',
    'stimulant_use',
    'inhalant_use',
    'heroin_use',
    'tranquilizer_use',
    'amphetamine_use'
]
correlation_matrix = data[mc_cols].corr()


vif_data = pd.DataFrame()
vif_data['feature_name'] = data[mc_cols].columns
vif_data['vif_score'] = [variance_inflation_factor(data[mc_cols].values, i) for i in range(len(mc_cols))]

print("\nVariance Inflation Factor (VIF):\n", vif_data)
```



```
Variance Inflation Factor (VIF):
      feature_name  vif_score
0  psychological_distress  1.180237
1           marijuana_use  1.182923
2       hallucinogen_use  1.028839
3           cocaine_use  1.094480
4           opioid_use  1.071641
5       stimulant_use  1.020335
6       inhalant_use  1.014413
7           heroine_use  1.137491
8   tranquilizer_use  1.067827
9       amphetamine_use  1.074707
```

```
In [24]: plt.figure(figsize= (8, 6))
sns.heatmap(
    correlation_matrix,
    annot= True,
    cmap= 'coolwarm',
    fmt= '.2f',
    linewidths= 0.5,
    square= True
)
plt.show()
```



Fortunately all features showed negligible correlation rates and low VIF scores, meaning all features are independently and collectively viable influencers to the analysis.

Prediction Models: Health Level

Challenge: `Health Level` is a respondent-percieved status on their overall health disposition, with 1-5 representing bad to great. Can this be directly imputed from given data? Can it be imputed specifically from drug use and addiction statuses?

Full Featureset

```
In [25]: # preprocess transformers
tf_minmax = MinMaxScaler()
tf_onehot = OneHotEncoder()
tf_ordinal = OrdinalEncoder()
tf_poweryeo = PowerTransformer(method= 'yeo-johnson')
tf_quantnorm = QuantileTransformer(output_distribution= 'normal')
tf_quantuni = QuantileTransformer(output_distribution= 'uniform')

# pipeline
preprocessor = ColumnTransformer(
    transformers=[
        ('tf1', tf_minmax, cols_minmax),
        ('tf2', tf_onehot, cols_onehot),
        ('tf3', tf_ordinal, cols_ordinal),
        ('tf5', tf_poweryeo, cols_poweryeo),
        ('tf6', tf_quantnorm, [x for x in cols_quantnorm if x != 'health_level']),
        ('tf7', tf_quantuni, cols_quantuni)
    ]
)

# datasplit
X_train, X_test, y_train, y_test = train_test_split(
    data_clean.drop(columns= 'health_level'),
    data_clean['health_level'],
    train_size= .2,
    random_state= 13
)

# apply fit transformations
X_train_scaled = preprocessor.fit_transform(X_train)
X_test_scaled = preprocessor.transform(X_test)
```

Logistic Classifier

```
In [26]: model = LogisticRegression(random_state= 13, max_iter= 1000)
model.fit(X_train_scaled, y_train)

y_pred = model.predict(X_test_scaled)
y_pred_train = model.predict(X_train_scaled)
```

```

acc = accuracy_score(y_train, y_pred_train)

print(f"Overall Train Accuracy: {acc:.4f}")

acc = accuracy_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)

print(f"Overall Test Accuracy: {acc:.4f}")
print(f"Mean Absolute Error: {mae:.4f}")

```

Overall Train Accuracy: 0.4575

Overall Test Accuracy: 0.4477

Mean Absolute Error: 0.6364

Logistic Forward Selection

```

In [27]: model = LogisticRegression(max_iter= 500, random_state= 13)
sfs = SequentialFeatureSelector(
    model,
    n_features_to_select= 'auto',
    direction= 'forward',
    cv= 5
)
sfs.fit(X_train_scaled, y_train)

selected_features = sfs.get_support(indices= True)
X_train_selected = X_train_scaled[:, selected_features]
X_test_selected = X_test_scaled[:, selected_features]
model.fit(X_train_selected, y_train)

y_pred = model.predict(X_test_selected)
y_pred_train = model.predict(X_train_selected)

acc = accuracy_score(y_train, y_pred_train)

print(f"Overall Train Accuracy: {acc:.4f}")

acc = accuracy_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)

print(f"Overall Test Accuracy: {acc:.4f}")
print(f"Mean Absolute Error: {mae:.4f}")
print("\nClassification Report:\n", classification_report(y_test, y_pred))

```

Overall Train Accuracy: 0.4566

Overall Test Accuracy: 0.4481

Mean Absolute Error: 0.6342

Classification Report:

	precision	recall	f1-score	support
1.0	0.17	0.00	0.01	229
2.0	0.20	0.00	0.00	2175
3.0	0.44	0.45	0.45	8068
4.0	0.45	0.73	0.56	9792
5.0	0.46	0.11	0.18	5083
accuracy			0.45	25347
macro avg	0.34	0.26	0.24	25347
weighted avg	0.43	0.45	0.39	25347

```
In [28]: preprocessor.get_feature_names_out()[selected_features]
```

```
Out[28]: array(['tf1_bmi', 'tf2_education_flag_False',  
                'tf2_education_flag_True', 'tf2_employment_Unemployed',  
                'tf2_marital_Divorced', 'tf2_marital_Married',  
                'tf2_nicotine_dependence_flag_False',  
                'tf2_nicotine_dependence_flag_True', 'tf2_sexuality_Homosexual',  
                'tf3_generation', 'tf5_amphetamine_use', 'tf5_crack_use',  
                'tf5_hallucinogen_use', 'tf5_heroin_use', 'tf5_inhalant_use',  
                'tf5_marijuana_use', 'tf5_psychological_distress',  
                'tf5_sedative_use', 'tf5_stimulant_use', 'tf6_education_level'],  
               dtype=object)
```

Random Forest

```
In [29]: model = RandomForestRegressor(n_estimators= 50, random_state= 13)  
model.fit(X_train_scaled, y_train)  
  
y_pred = model.predict(X_test_scaled)  
y_pred_train = model.predict(X_train_scaled)  
  
r2 = r2_score(y_train, y_pred_train)  
  
print(f"Train R² Score: {r2:.4f}")  
  
mae = mean_absolute_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)  
  
print(f"Test R² Score: {r2:.4f}")  
print(f"Mean Absolute Error: {mae:.4f}")
```

Train R² Score: 0.8710

Test R² Score: 0.1186

Mean Absolute Error: 0.6909

Gradient Booster

```
In [30]: model = GradientBoostingRegressor(n_estimators= 50, learning_rate=0.1, max_depth=3,
model.fit(X_train_scaled, y_train)
y_pred = model.predict(X_test_scaled)

r2 = r2_score(y_train, y_pred_train)
print(f"Train R² Score: {r2:.4f}")

r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)

print(f"Test R² Score: {r2:.4f}")
print(f"Mean Squared Error: {mse:.4f}")
```

Train R² Score: 0.8710

Test R² Score: 0.1879

Mean Squared Error: 0.6848

Drug Use Sub-Featureset

```
In [31]: drug_subset = [
    'mental_illness_flag',
    'nicotine_dependence_flag',
    'amphetamine_use',
    'cocaine_use',
    'crack_use' ,
    'hallucinogen_use',
    'heroin_use',
    'inhalant_use',
    'marijuana_use',
    'opioid_use',
    'sedative_use',
    'stimulant_use',
    'tranquilizer_use',
    'health_level'
]
```

```
In [32]: data_clean_subset1 = data_clean[drug_subset]
tf_onehot = OneHotEncoder(drop= 'first')
tf_poweryeo = PowerTransformer(method= 'yeo-johnson')

preprocessor = ColumnTransformer(
    transformers=[
        ('tf2', tf_onehot, [
            'mental_illness_flag',
            'nicotine_dependence_flag'
        ]),
        ('tf5', tf_poweryeo, [
            'amphetamine_use',
            'cocaine_use',
            'crack_use' ,
            'hallucinogen_use',
            'heroin_use',
            'inhalant_use',
            'marijuana_use',
```

```

        'opioid_use',
        'sedative_use',
        'stimulant_use',
        'tranquilizer_use'
    ])
]
)

X_train, X_test, y_train, y_test = train_test_split(
    data_clean_subset1.drop(columns= 'health_level'),
    data_clean_subset1['health_level'],
    train_size= .2,
    random_state= 13
)

X_train_scaled = preprocessor.fit_transform(X_train)
X_test_scaled = preprocessor.transform(X_test)

```

Logistic Classifier

```

In [33]: model = LogisticRegression(random_state= 13, max_iter= 1000)
model.fit(X_train_scaled, y_train)
y_pred = model.predict(X_test_scaled)
y_pred_train = model.predict(X_train_scaled)

acc = accuracy_score(y_train, y_pred_train)

print(f"Overall Train Accuracy: {acc:.4f}")

acc = accuracy_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)

print(f"Overall Test Accuracy: {acc:.4f}")
print(f"Mean Absolute Error: {mae:.4f}")

```

Overall Train Accuracy: 0.4217

Overall Test Accuracy: 0.4086

Mean Absolute Error: 0.6752

Generation Predictor

```

In [34]: generation_subset = [
    'generation',
    'mental_illness_flag',
    'nicotine_dependence_flag',
    'amphetamine_use',
    'cocaine_use',
    'crack_use',
    'hallucinogen_use',
    'heroin_use',
    'inhalant_use',
    'marijuana_use',
    'opioid_use',
    'sedative_use',
    'stimulant_use',

```

```
    'tranquilizer_use'  
]
```

Random Forest

```
In [35]: data_clean_subset2 = data_clean[generation_subset].copy()  
tf_poweryeo = PowerTransformer(method= 'yeo-johnson')  
  
label_encoder = LabelEncoder()  
data_clean_subset2['generation'] = \  
    label_encoder.fit_transform(data_clean_subset2['generation'])  
data_clean_subset2['mental_illness_flag'] = \  
    data_clean_subset2['mental_illness_flag'].map({'True': 1, 'False': 0})  
data_clean_subset2['nicotine_dependence_flag'] = \  
    data_clean_subset2['nicotine_dependence_flag'].map({'True': 1, 'False': 0})  
data_clean_subset2[[  
    'amphetamine_use',  
    'cocaine_use',  
    'crack_use' ,  
    'hallucinogen_use',  
    'heroin_use',  
    'inhalant_use',  
    'marijuana_use',  
    'opioid_use',  
    'sedative_use',  
    'stimulant_use',  
    'tranquilizer_use'  
]] = tf_poweryeo.fit_transform(data_clean_subset2[[  
    'amphetamine_use',  
    'cocaine_use',  
    'crack_use' ,  
    'hallucinogen_use',  
    'heroin_use',  
    'inhalant_use',  
    'marijuana_use',  
    'opioid_use',  
    'sedative_use',  
    'stimulant_use',  
    'tranquilizer_use'  
]])  
  
X = data_clean_subset2.drop(columns=['generation'])  
y = data_clean_subset2['generation']  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size= 0.2, random_st  
model = RandomForestClassifier(n_estimators= 50, random_state= 13)  
  
model.fit(X_train, y_train)  
  
y_pred = model.predict(X_test)  
y_pred_train = model.predict(X_train)  
  
r2 = accuracy_score(y_train, y_pred_train)  
print(f"Overall Train Accuracy: {r2:.4f}")  
  
r2 = accuracy_score(y_test, y_pred)
```

```
mae = mean_absolute_error(y_test, y_pred)

print(f"Overall Test Accuracy: {r2:.4f}")
print(f"Mean Absolute Error: {mae:.4f}")
```

Overall Train Accuracy: 0.3743
 Overall Test Accuracy: 0.3618
 Mean Absolute Error: 1.0642

Support Vector Machines

```
In [36]: model = SVC(kernel='rbf', C= 1.0, gamma= 'scale', random_state= 13)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
y_pred_train = model.predict(X_train)

acc = accuracy_score(y_train, y_pred_train)
print(f"Overall Train Accuracy: {acc:.4f}")

acc = accuracy_score(y_test, y_pred)
mae = mean_absolute_error (y_test, y_pred)

cr = classification_report(y_test, y_pred, target_names= label_encoder.classes_, ze
print(f"Overall Test Accuracy: {acc:.4f}")
print(f"Mean Absolute Error: {mae:.4f}")
print(f"\nClassification Report:\n{cr}")
```

Overall Train Accuracy: 0.3649
 Overall Test Accuracy: 0.3617
 Mean Absolute Error: 1.0704

Classification Report:

	precision	recall	f1-score	support
Baby Boomer	1.00	0.00	0.00	208
Millennial	0.22	0.02	0.03	1437
X	0.00	0.00	0.00	648
Xennial	0.35	0.67	0.46	2024
Zillennial	0.39	0.45	0.42	2020
accuracy			0.36	6337
macro avg	0.39	0.23	0.18	6337
weighted avg	0.32	0.36	0.29	6337

Conclusion

PCA was successful in reducing the dimensions of our data, with explained variances in excess of 90% distribution. Component vectors showed the drug usage features have the highest consistent influence. This is an expected and positive result because it shows consistent viability of the drug usage data in classifying survey respondents. Integrity validations will be needed to ensure there is no dependency factor within these features

because the usage of them is not assuredly mutually exclusive. Multicollinearity analysis later showed these dependencies were not an issue, supporting the integrity of the PCA results.

Overall, even though PCA was successful, analysis around the drug use data yielded unusable predictive models. Health Levels were modeled with subpar results and accuracies ranging between sub-random chance (20%) and 50%. Efforts to identify generational classifications were even less successful.

The mixed results of this report show poor modeling efficacy on the entire collection of respondents. Post-analysis research on other peer journals with the survey data show different methodologies on focused demographic subsets had more informative results. Most analysis focused on adolescents, gender, and the rise in opioid use as time-series analysis over several survey report years;^[1] ^[2] other reports were able to expand and compare the survey responses to comparable clinical analysis.^[3]

^1. National Estimates of Marijuana Use and Related Indicators
<http://dx.doi.org/10.15585/mmwr.ss6511a1>

^2. Gender and prescription opioids: Findings from the National Survey on Drug Use and Health <https://doi.org/10.1016/j.addbeh.2010.06.018>

^3. A clinical validation of the National Survey on Drug Use and Health Assessment of Substance Use Disorders <https://doi.org/10.1016/j.addbeh.2007.12.007>